

The Model Lake

3,016,439 model embeddings, 18,729 dataset–task queries, and the two-stage retrieval that fishes one from the other

3,016,439

candidates

128-D unit sphere
one shared retrieval space

1,000

pooled

HNSW inner product
recall@1000 0.9923 · p50 0.406 ms

10

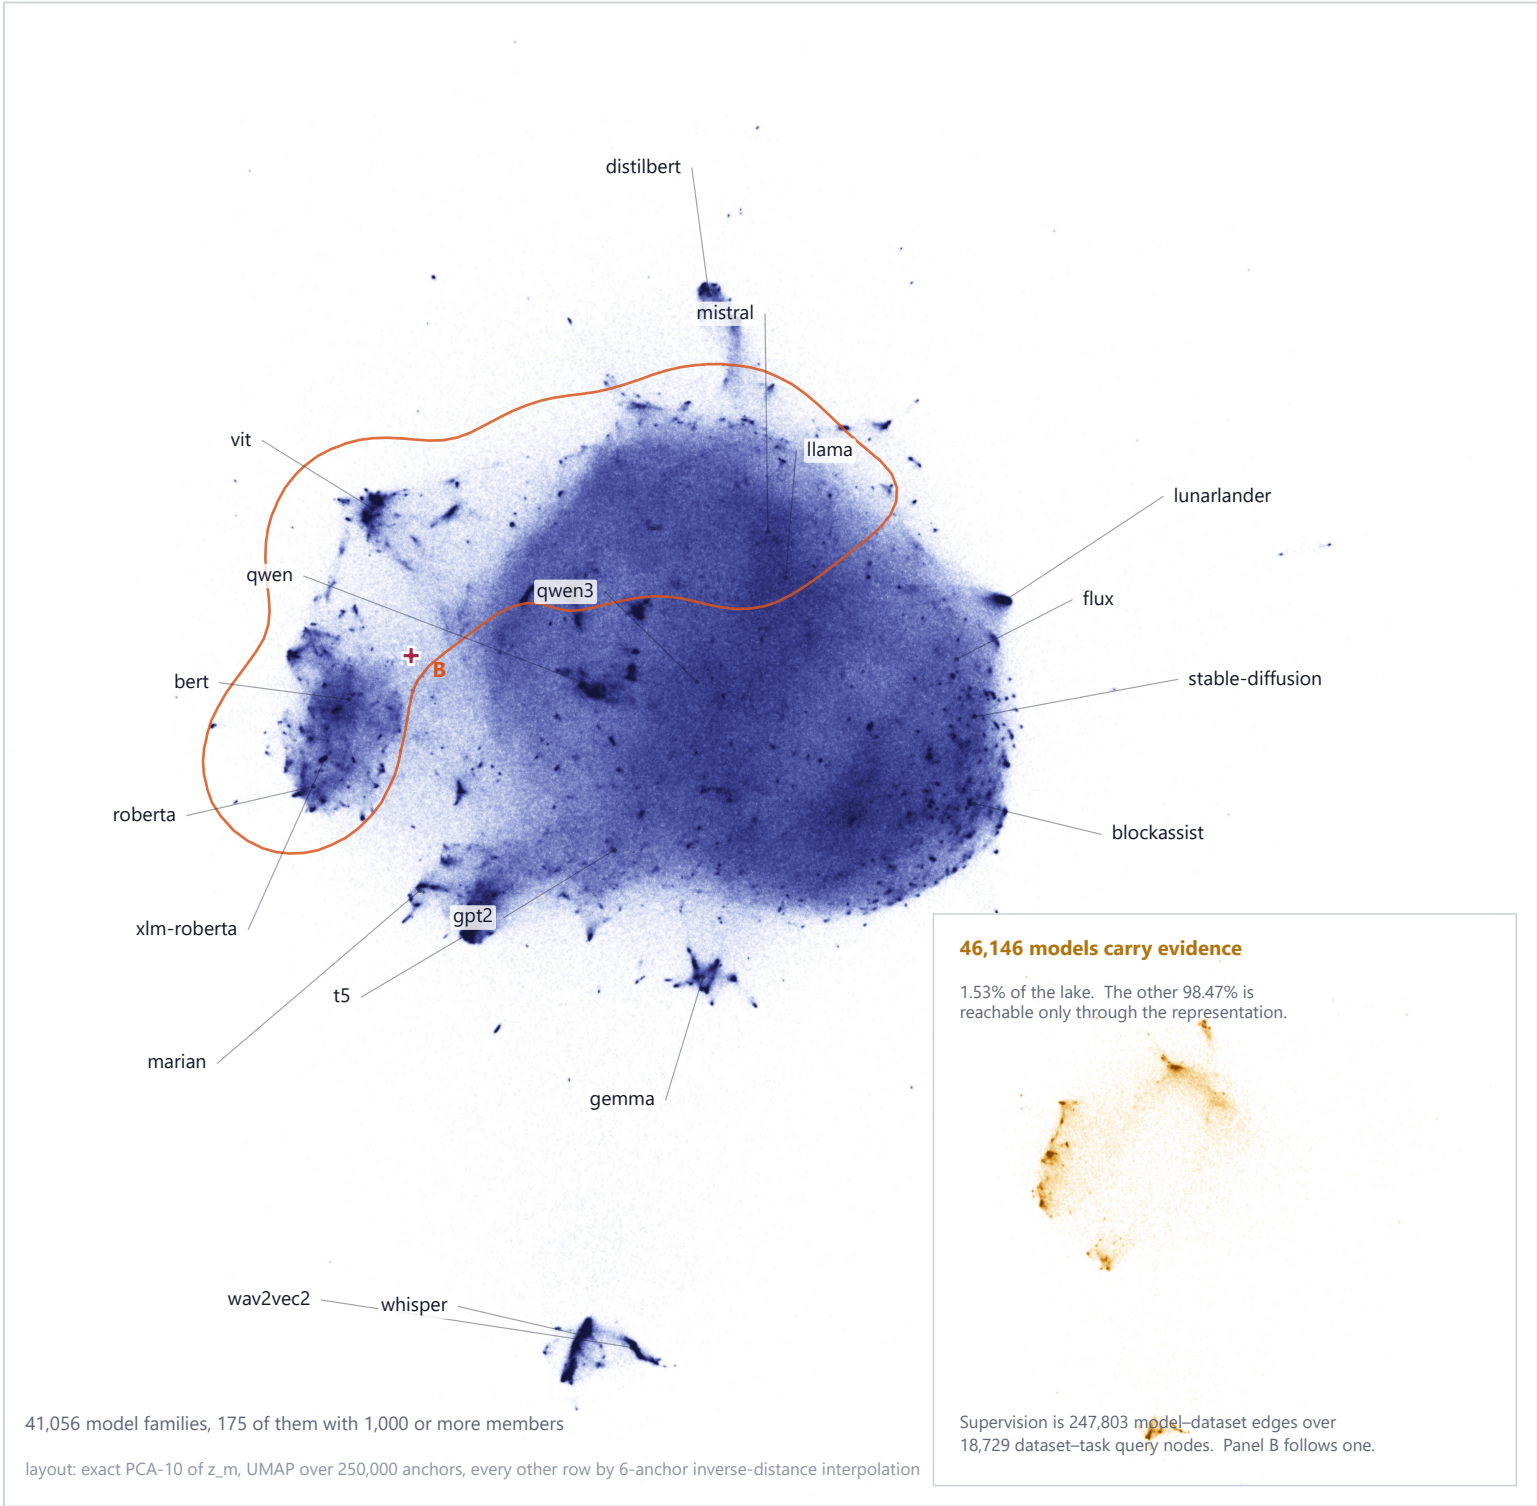
returned

task-prior rerank
p50 0.129 ms · gold@10 0.3279

seed 0 throughout; three-seed mean gold@10 0.3031

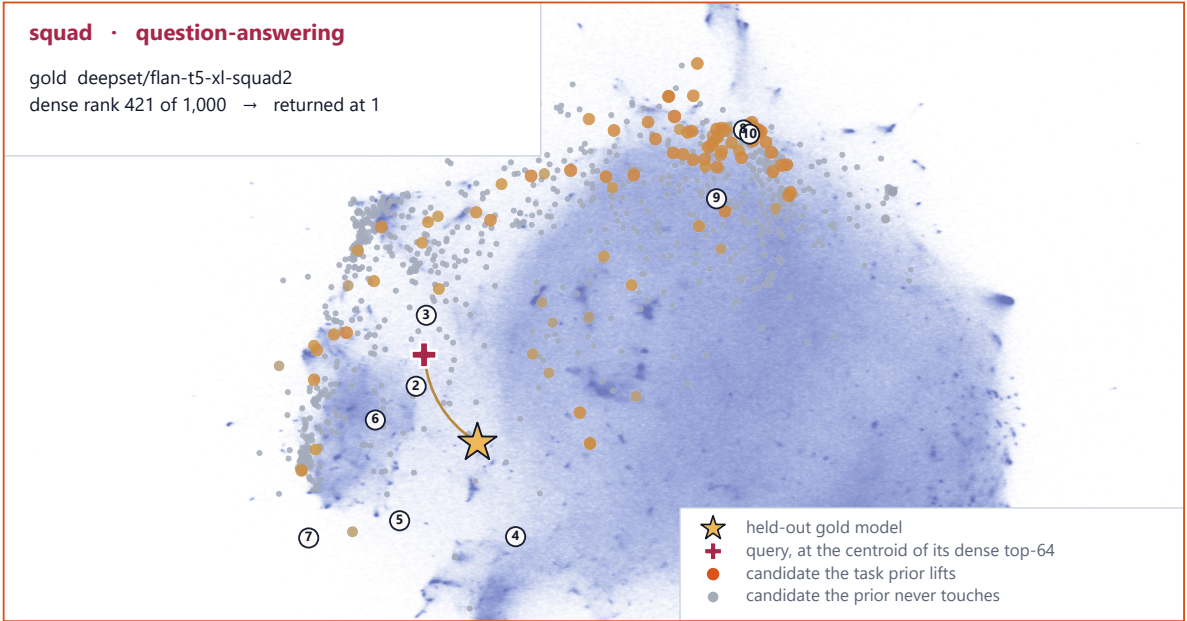
A the lake, laid out once

3,016,439 model rows, 410 of them beyond this frame · log density · neighbourhoods are meaningful, distances are not



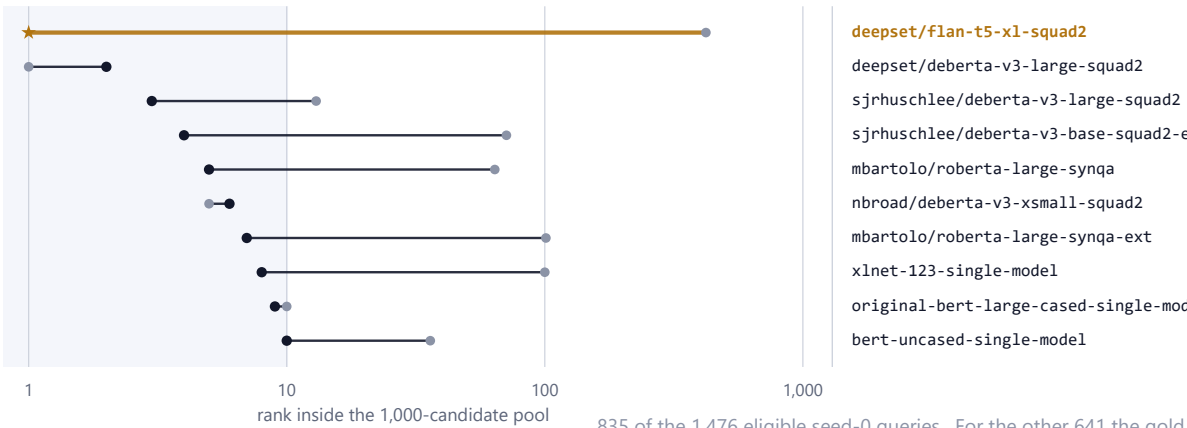
B one cast

all 1,000 pooled candidates, enlarged from the outline in A



C what came back

grey dot: where the dense stage had it · bar: how far the prior moved it



D rank migration

